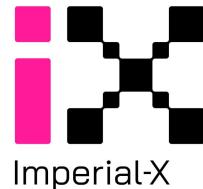


# Deep Learning in Python

Lecture 2: Convolutional Neural Networks

Dr Andreas C. S. Jørgensen

Imperial College  
London



# Outline

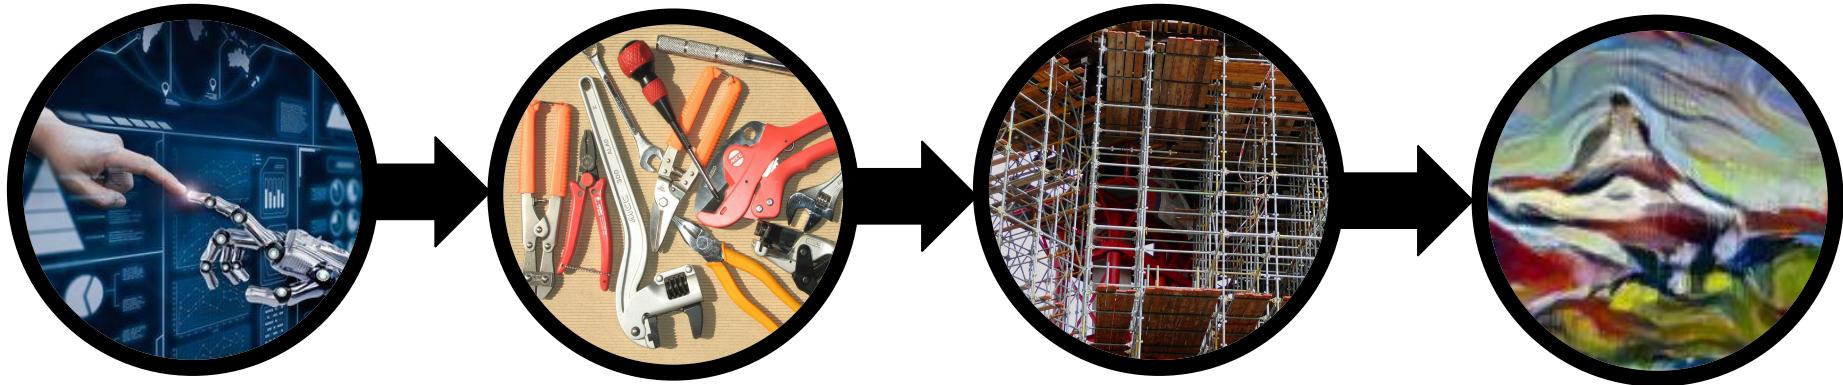
**Hands-on course** in Deep Learning using Python (PyTorch) with focus on computer vision.

<b>Lecture 1</b>	Deep Learning concepts and terminology. PyTorch commands, syntax and libraries (torch, torchvision). Building deep neural networks.
<b>Lecture 2</b>	Convolutional Neural Networks (CNN). Building CNN in PyTorch. Using GPUs with PyTorch.
<b>Lecture 3</b>	Optimisation: How to train and test neural networks robustly. Regularisation and hyper-parameter tuning.

# Classify images



# Outline



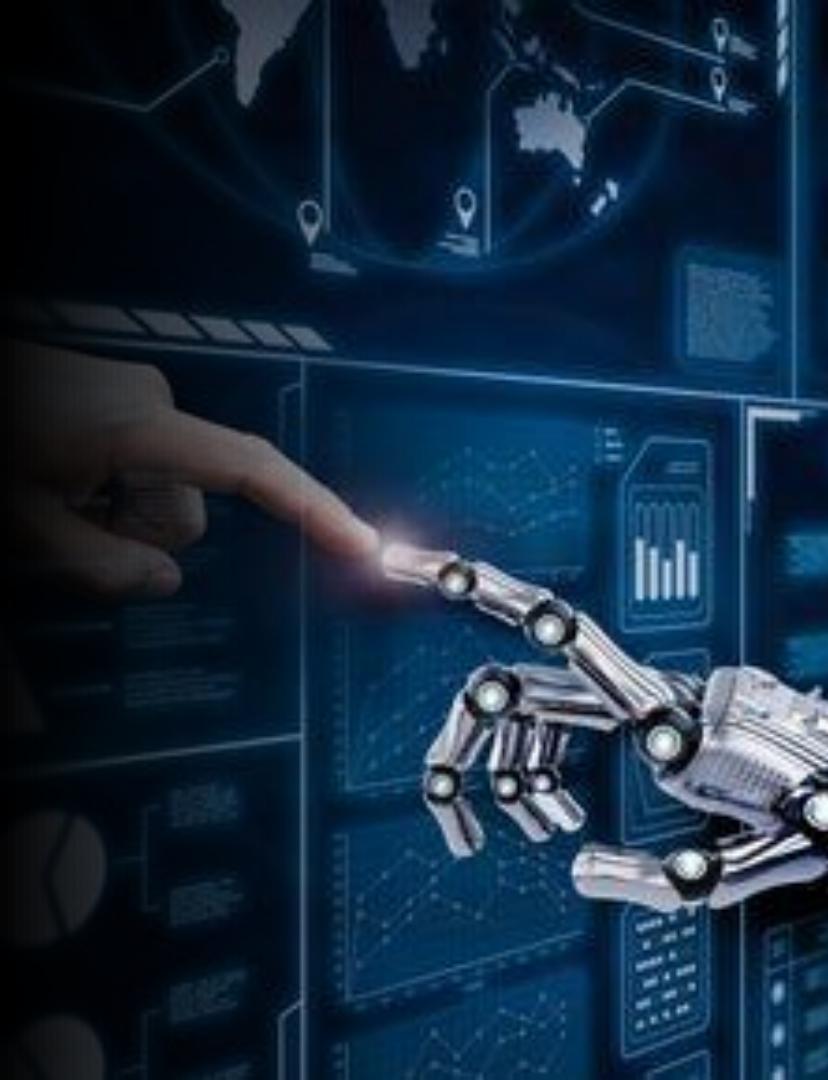
1) Concepts

2) Toolkit (PyTorch)

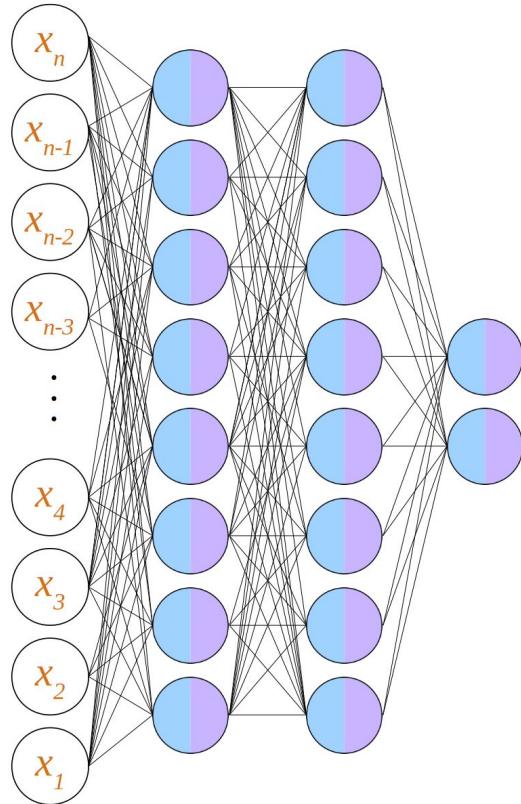
3) Combine 1 and 2  
to build CNN

4) Beyond  
classification

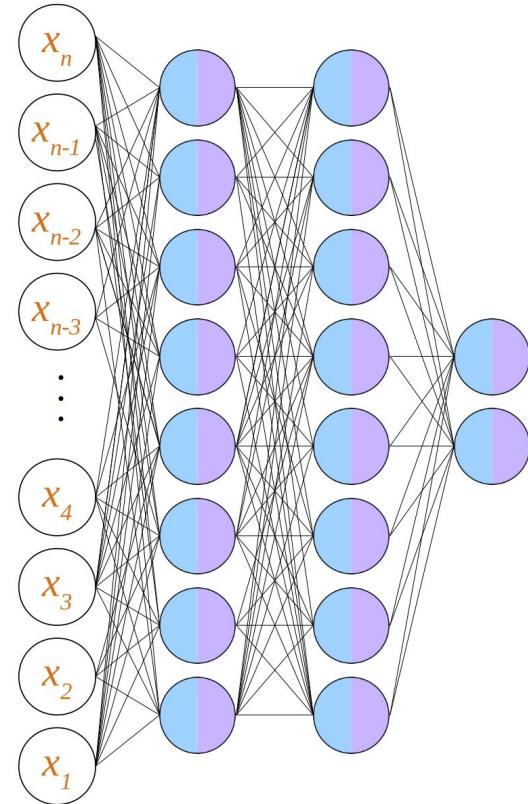
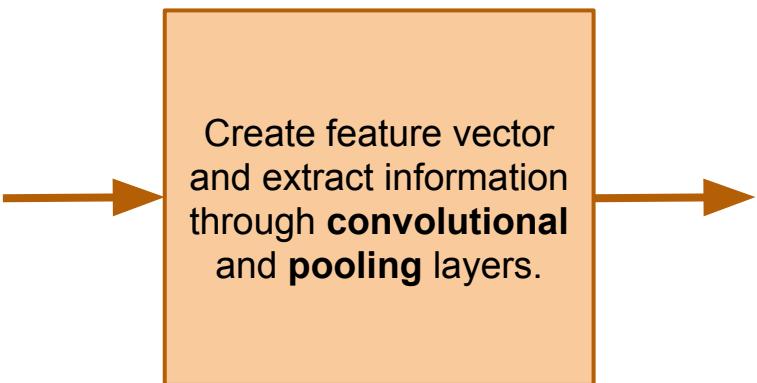
# Convolutional Neural Networks



# Architecture



# Architecture



# Convolutional layers



2	5	8		7	9	8
1	3	4		2	8	3
0	4	1		2	1	4
3	3	8		7	5	1
4	1	2		2	9	2
5	1	3		7	4	8

Image

-1	-2	-1
0	0	0
1	2	1

\* Kernel/Filter  
(Horizontal edge detection)

-11			

Resulting image

**Convolution:** Sum result of element-wise multiplication with kernel

# Convolutional layers



2	5	8	7	9	8
1	3	4	2	8	3
0	4	1	2	1	4
3	3	8	7	5	1
4	1	2	2	9	2
5	1	3	7	4	8

Image

-1	-2	-1
0	0	0
1	2	1

\* Kernel/Filter  
(Horizontal edge detection)

-11	-20		

Resulting image

**Stride:** The kernel moves (strides) a number of pixels for every convolution operation

# Convolutional layers



2	5	8	7	9	8
1	3	4	2	8	3
0	4	1	2	1	4
3	3	8	7	5	1
4	1	2	2	9	2
5	1	3	7	4	8

Image

-1	-2	-1
0	0	0
1	2	1

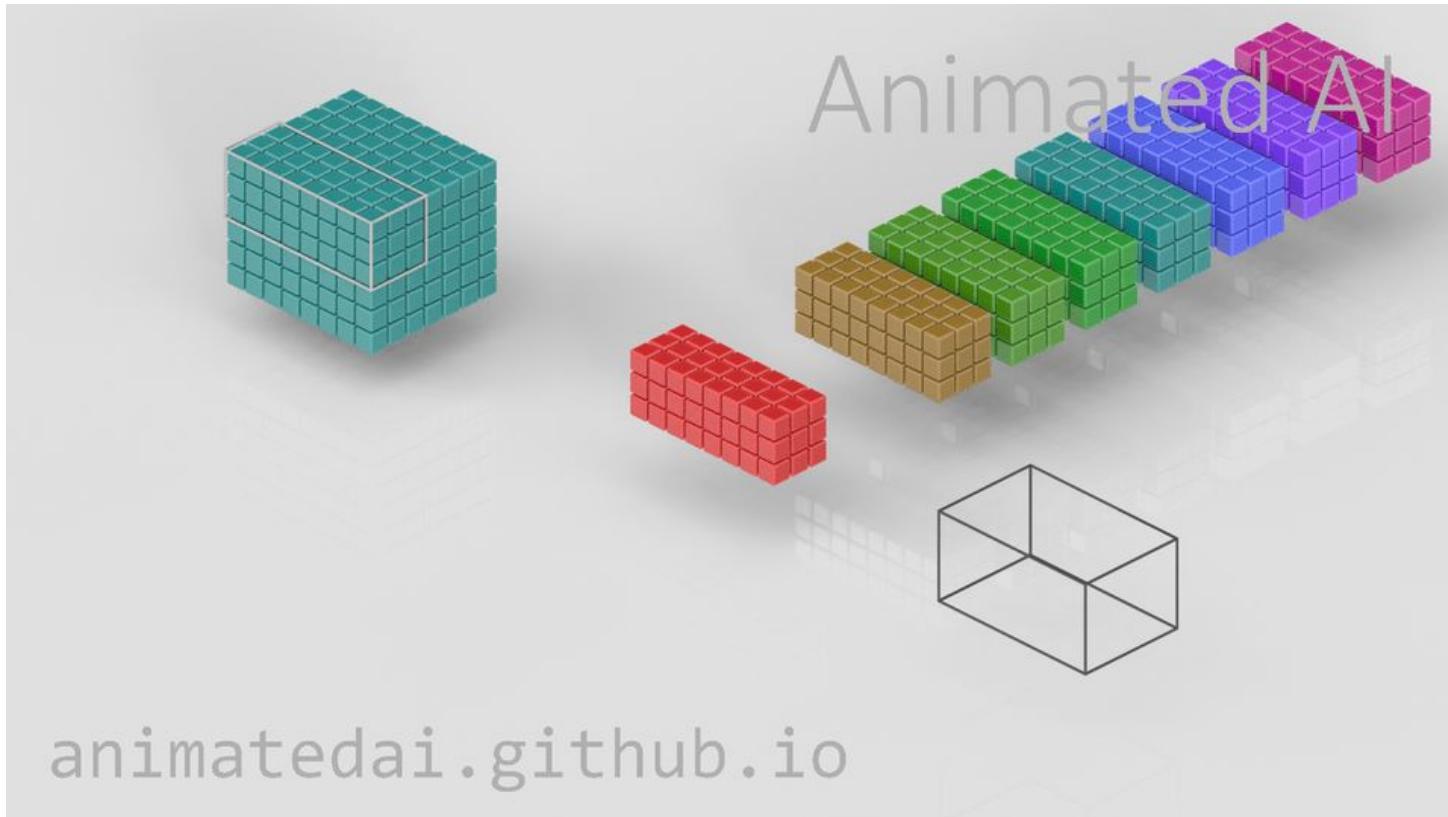
\* Kernel/Filter  
(Horizontal edge detection)

-11	-20	-45	-70
6	13	11	-3
-1	-1	9	14
-7	-12	-6	5

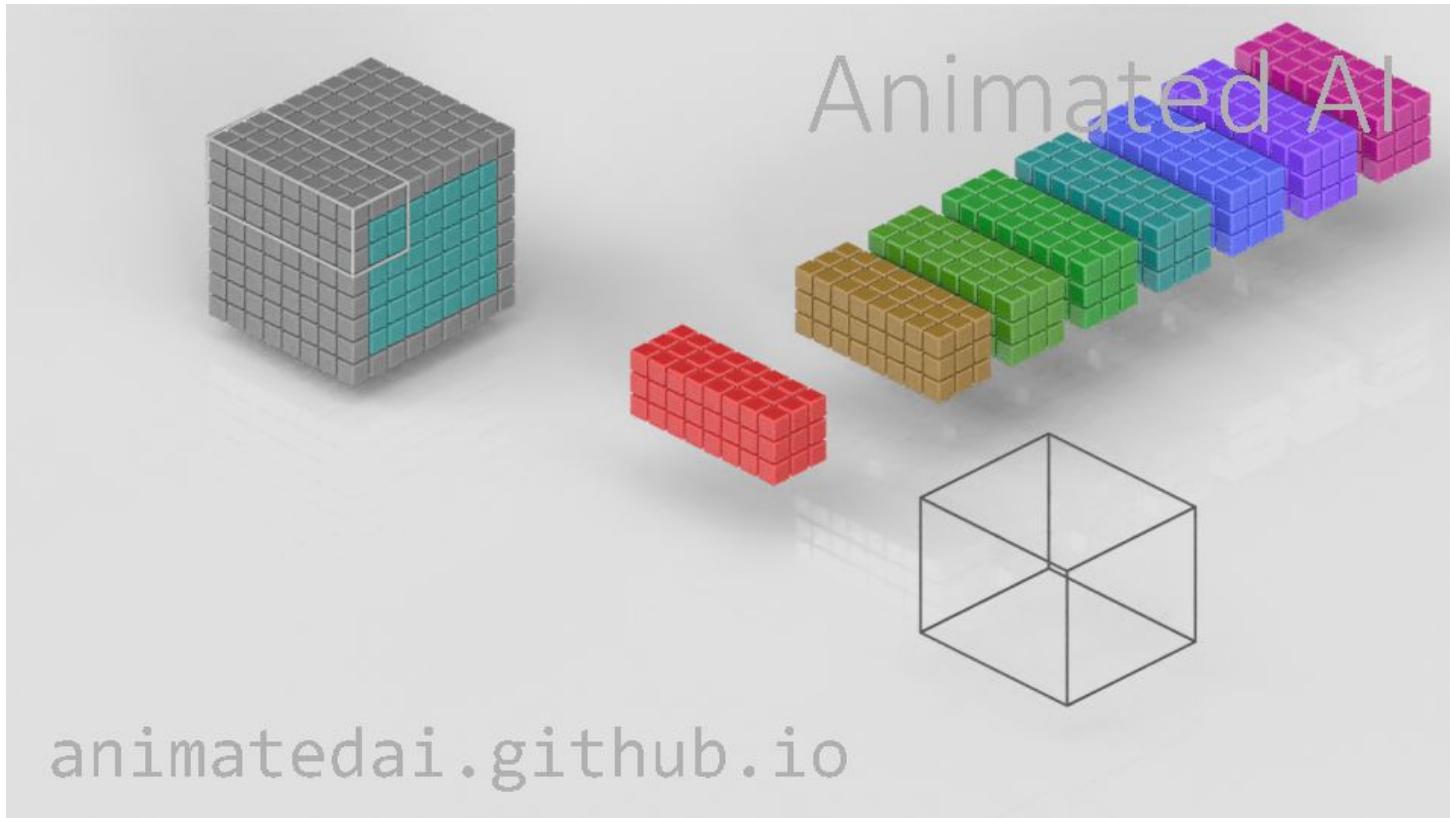
Resulting image

**Stride:** The kernel moves (strides) a number of pixels for every convolution operation

# Convolutional layers



# Padding



# Other convolutions



- Other dimensions than  $3 \times 3$ .
- Depthwise convolutions: One filter per channel. The output is stacked.
- $1 \times 1$  (pointwise) convolutions: Add up channels.
- Dilated convolutions: Filters with gaps (gives a wider field of view).
- Transposed convolutions: Increase spatial dimensions adding details.

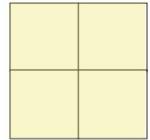
# Pooling layers



2	5	8	7	9	8
1	3	4	2	8	3
0	4	1	2	1	4
3	3	8	7	5	1
4	1	2	2	9	2
5	1	3	7	4	8

Image

\*



Pooling window

=

5		

Resulting image

**Pooling:** Reduce feature map size and add spatial invariance.

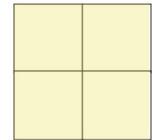
# Pooling layers



2	5	8	7	9	8
1	3	4	2	8	3
0	4	1	2	1	4
3	3	8	7	5	1
4	1	2	2	9	2
5	1	3	7	4	8

Image

\*



=

5	8	

\* Pooling window  
(Max pooling)

Resulting image

**2×2 max pooling:** Finds maximum value within window: The most prominent feature. With a stride of 2.

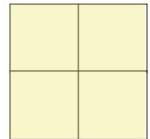
# Pooling layers



2	5	8	7	9	8
1	3	4	2	8	3
0	4	1	2	1	4
3	3	8	7	5	1
4	1	2	2	9	2
5	1	3	7	4	8

Image

\*



=

5	8	9
4	8	5
5	7	9

Resulting image

\* Pooling window  
(Max pooling)

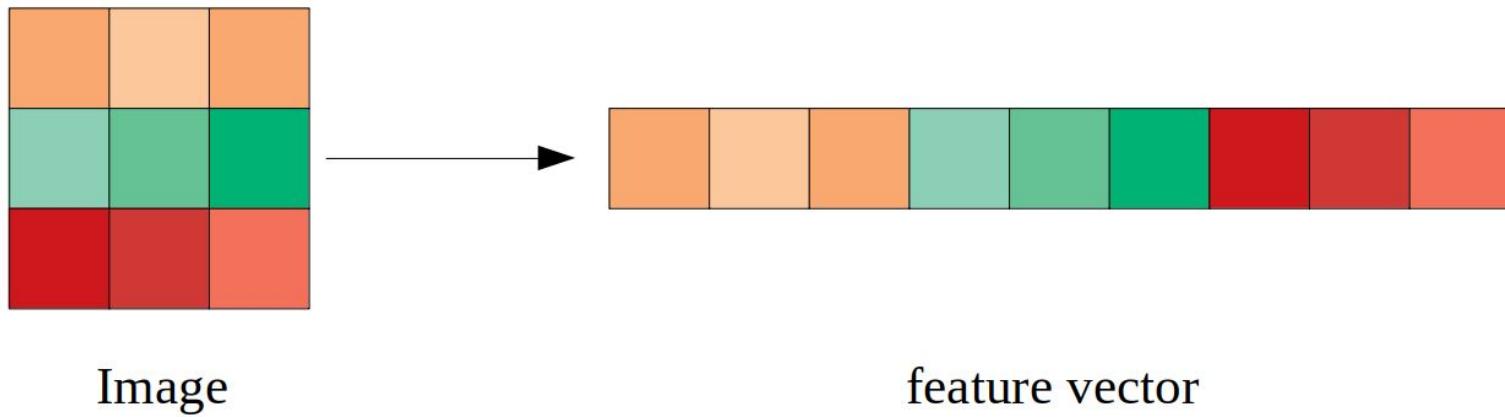
**2×2 max pooling:** Finds maximum value within window: The most prominent feature. With a stride of 2.

# Pooling layers



- $2 \times 2$  or  $3 \times 3$  max pooling: Retains most prominent feature. Even when moved by a pixel, we get the same result, i.e., the pooling is not sensitive to position.
- Average pooling: Smoothed representation.
- Global average pooling: The global average of the entire picture.
- Adaptive pooling: Vary the pooling stride and window size.

# Flattening



We now have a feature vector and can feed the data into a neural network.

# Transfer learning



Rather than training a new network, we just *tweak* an existing one. First, we replace the output layer with a layer that produces the right number of outputs.

**Fine tuning:** Adjust the parameters through training on the data set.

**Feature extraction:** Only calibrate the final layer, but "freeze" the rest of the network.

Only viable option for **small datasets!**

# Step by step



Components of a Convolutional Neural Network:

- Use convolutional layers to detect patterns and features (e.g., edges, textures).
- Use pooling layers to reduce feature map size and introduce spatial invariance.
- Pass the extracted features to the type of architecture we have seen: Fully-connected layers and an output layer (using the cross-validation loss for classification).

# New tools required?



# GPUs



Graphics processing units can perform many computations in parallel and can thus lead to speed up for parallelizable tasks.

*# Check GPU availability*

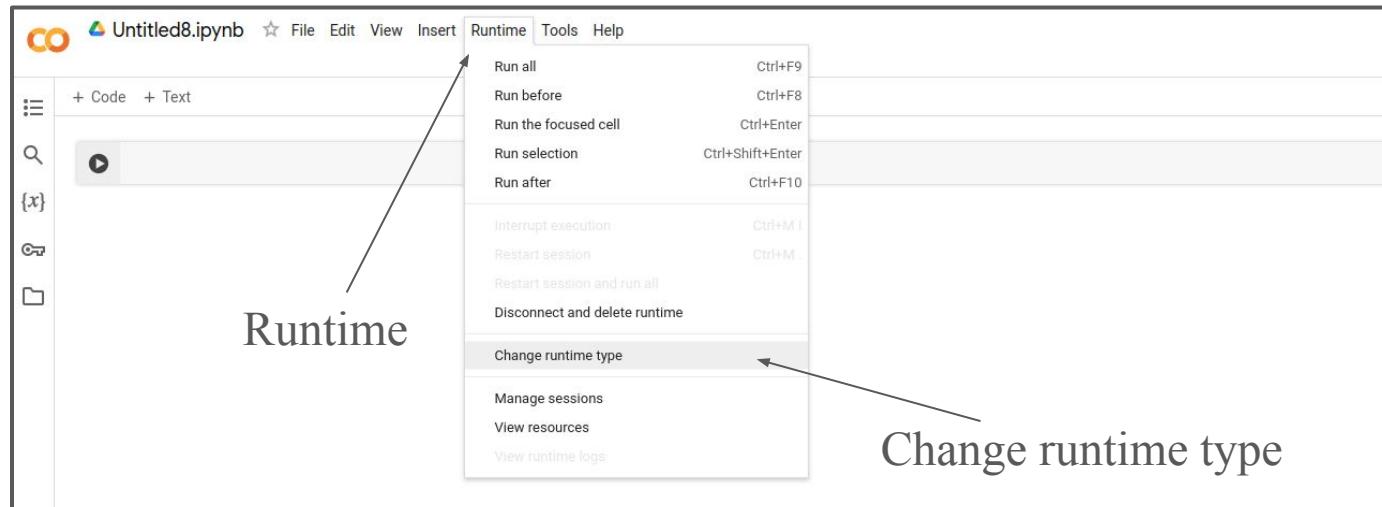
```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

# GPUs

Graphics processing units can perform many computations in parallel and can thus lead to speed up for parallelizable tasks.

```
# Check GPU availability
```

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

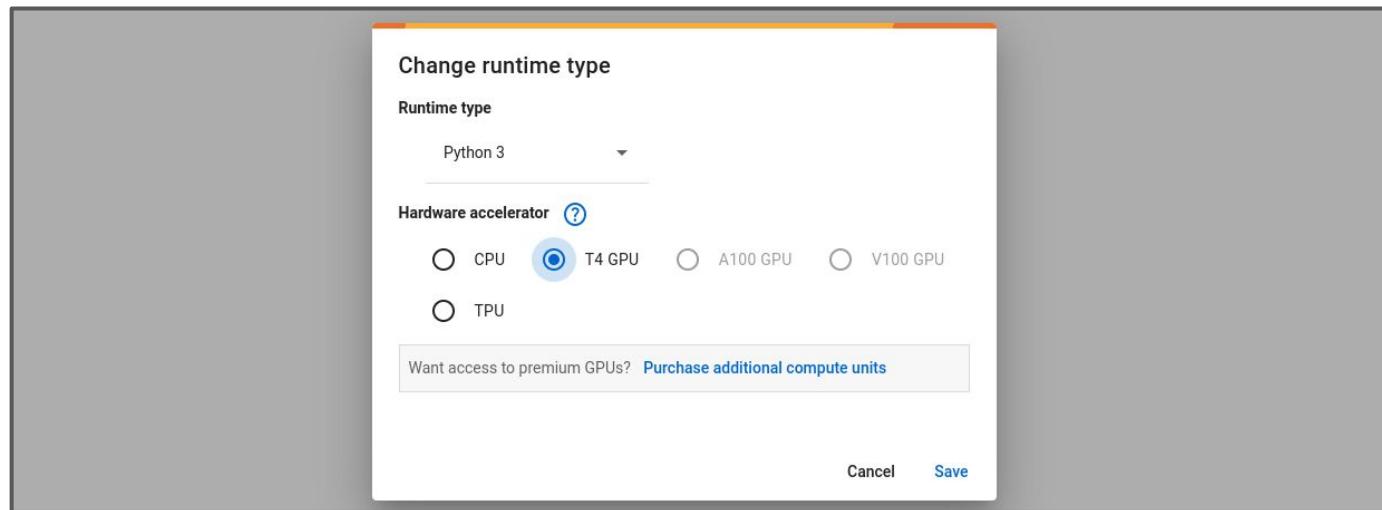


# GPUs

Graphics processing units can perform many computations in parallel and can thus lead to speed up for parallelizable tasks.

```
# Check GPU availability
```

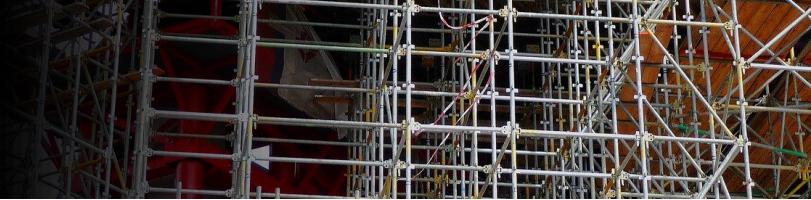
```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```



# Building CNNs



# New layers



```
class CNN(nn.Module):  
    def __init__(self):  
        super().__init__() # Inherit from parent class
```

*# Define layers and activations functions*

```
def forward(self, x):
```

*# Transform x stepwise, layer by layer:*

```
    return x
```

$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers



```
class CNN(nn.Module):
    def __init__(self):
        super().__init__() # Inherit from parent class

        self.act1 = nn.ReLU() # Activation function

        self.fc1 = nn.Linear(6*8*8, 120) # Fully connected
        self.fc2 = nn.Linear(120, 84) # Fully connected
        self.fc3 = nn.Linear(84, 10) # Fully connected

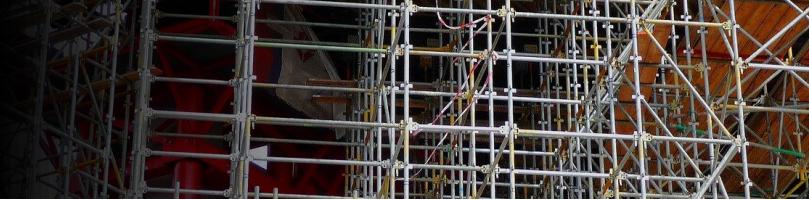
    def forward(self, x):

        # Transform x stepwise, layer by layer:

        return x
```

$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers

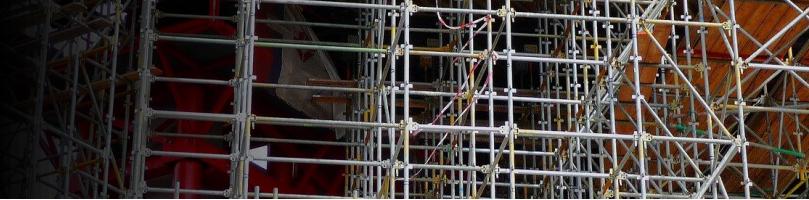


```
class CNN(nn.Module):
    def __init__(self):
        super().__init__() # Inherit from parent class
        self.conv1 = nn.Conv2d(3, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.act1 = nn.ReLU() # Activation function
        self.pool1 = nn.MaxPool2d(2, 2) # Pooling
        self.conv2 = nn.Conv2d(6, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.flat = nn.Flatten() # Flattening
        self.fc1 = nn.Linear(6*8*8, 120) # Fully connected
        self.fc2 = nn.Linear(120, 84) # Fully connected
        self.fc3 = nn.Linear(84, 10) # Fully connected

    def forward(self, x):
        # Transform x stepwise, layer by layer:
        return x
```

$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers

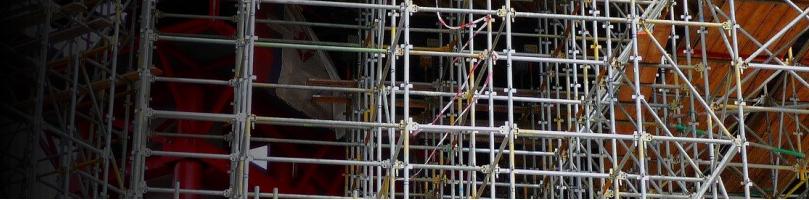


```
class CNN(nn.Module):
    def __init__(self):
        super().__init__() # Inherit from parent class
        self.conv1 = nn.Conv2d(3, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.act1 = nn.ReLU() # Activation function
        self.pool1 = nn.MaxPool2d(2, 2) # Pooling
        self.conv2 = nn.Conv2d(6, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.flat = nn.Flatten() # Flattening
        self.fc1 = nn.Linear(6*8*8, 120) # Fully connected   6 channels of 8x8 pixels (starting at 32x32 pixels)
        self.fc2 = nn.Linear(120, 84) # Fully connected
        self.fc3 = nn.Linear(84, 10) # Fully connected

    def forward(self, x):
        # Transform x stepwise, layer by layer:
        return x
```

$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers



```
class CNN(nn.Module):
    def __init__(self):
        super().__init__() # Inherit from parent class
        self.conv1 = nn.Conv2d(3, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.act1 = nn.ReLU() # Activation function
        self.pool1 = nn.MaxPool2d(2, 2) # Pooling
        self.conv2 = nn.Conv2d(6, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.flat = nn.Flatten() # Flattening
        self.fc1 = nn.Linear(6*8*8, 120) # Fully connected
        self.fc2 = nn.Linear(120, 84) # Fully connected
        self.fc3 = nn.Linear(84, 10) # Fully connected

    def forward(self, x):
        x = self.pool1(self.act1(self.conv2(self.pool1(self.act1(self.conv1(x)))))))
        x = self.flat(x)
        x = self.fc3(self.act1(self.fc2(self.act1(self.fc1(x))))))
        return x
```

$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers

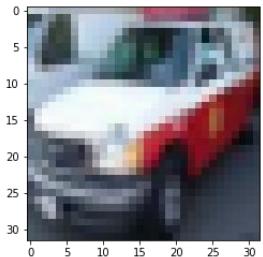
```
class CNN(nn.Module):
    def __init__(self):
        super().__init__() # Inherit from parent class
        self.conv1 = nn.Conv2d(3, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.act1 = nn.ReLU() # Activation function
        self.pool1 = nn.MaxPool2d(2, 2) # Pooling
        self.conv2 = nn.Conv2d(6, 6, kernel_size=(3,3), stride=1, padding=1) # Convolutional
        self.flat = nn.Flatten() # Flattening
        self.fc1 = nn.Linear(6*8*8, 120) # Fully connected
        self.fc2 = nn.Linear(120, 84) # Fully connected
        self.fc3 = nn.Linear(84, 10) # Fully connected

    def forward(self, x):
        x = self.pool1(self.act1(self.conv2(self.pool1(self.act1(self.conv1(x)))))))
        x = self.flat(x)
        x = self.fc3(self.act1(self.fc2(self.act1(self.fc1(x))))))
        return x
```

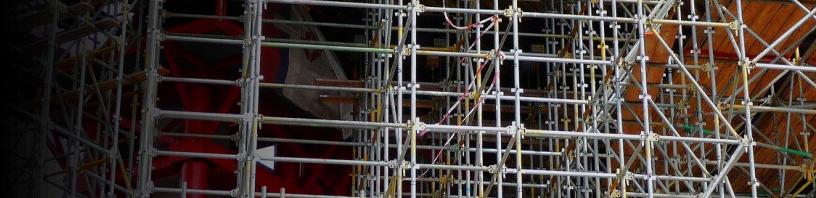
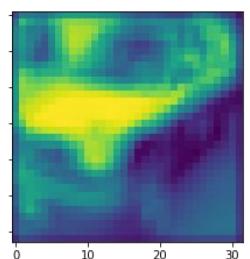
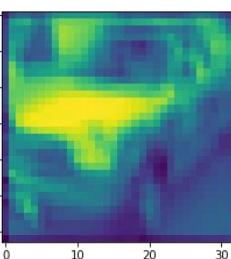
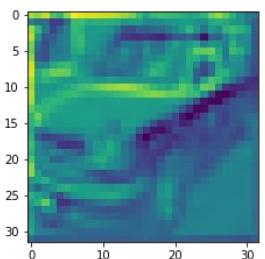
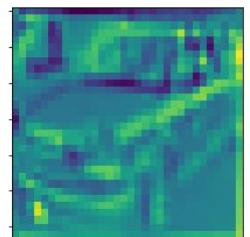
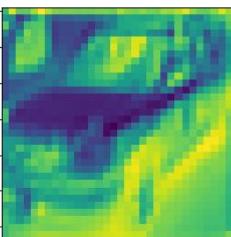
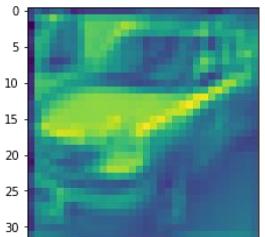
$$\Phi(\mathbf{x}) = \phi_K(\phi_{K-1}(\dots \phi_2(\phi_1(\mathbf{x})) \dots))$$

# New layers

```
X = torch.tensor([trainset.data[60]], dtype=torch.float32).permute(0,3,1,2)
with torch.no_grad():
    X = cnn.conv1(X)
```



$\text{conv1}(X)$



# Training

```
loss_func = nn.CrossEntropyLoss()  
optimizer = optim.SGD(cnn.parameters(), lr=1e-3, momentum=0.9)
```

# Training



```
loss_func = nn.CrossEntropyLoss()  
optimizer = optim.SGD(cnn.parameters(), lr=1e-3, momentum=0.9)
```

In each epoch, looping over batches,

- 1) Make model predictions.
- 2) Compute the loss.
- 3) Compute the gradient of the loss.
- 4) Update parameters.

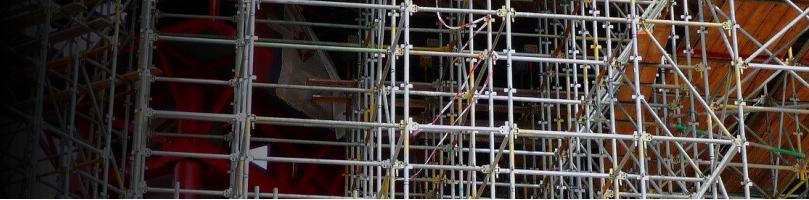
```
for epoch in range(n_epochs):  
    for data in trainloader:  
        inputs, labels = data  
        # 1) Feedforward  
        y_pred = cnn(inputs)  
  
        # 2) Compute the loss  
        loss = loss_func(y_pred, labels)  
  
        # 3) Backpropagation of errors  
        optimizer.zero_grad()  
        loss.backward()  
  
        # 4) Update parameters  
        optimizer.step()
```

# Transfer learning

Replace the output layer with a layer that produces the right number of outputs.

**Fine tuning:** Adjust the parameters through training on the data set.

# Transfer learning



Replace the output layer with a layer that produces the right number of outputs.

**Fine tuning:** Adjust the parameters through training on the data set.

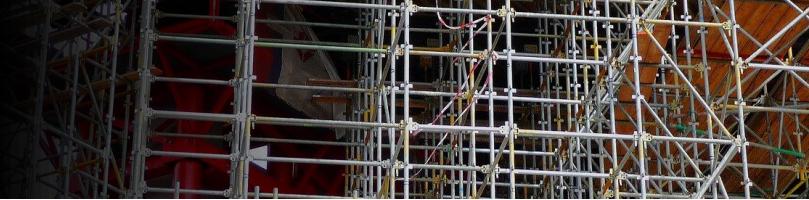
```
from torchvision.models import resnet50

# Instigate the pretrained model
model1 = resnet50(pretrained=True)

# Input features of the last layer of the original model
n_ftrs = model1.fc.in_features

# Replace the last layer with a new output layer that matches the dataset.
n_classes = len(class_names)
model1.fc = nn.Linear(n_ftrs, n_classes)
```

# Transfer learning



Replace the output layer with a layer that produces the right number of outputs.

**Fine tuning:** Adjust the parameters through training on the data set.

```
from torchvision.models import resnet50

# Instigate the pretrained model
model1 = resnet50(pretrained=True)

# Input features of the last layer of the original model
n_ftrs = model1.fc.in_features

# Replace the last layer with a new output layer that matches the dataset.
n_classes = len(class_names)
model1.fc = nn.Linear(n_ftrs, n_classes)
```

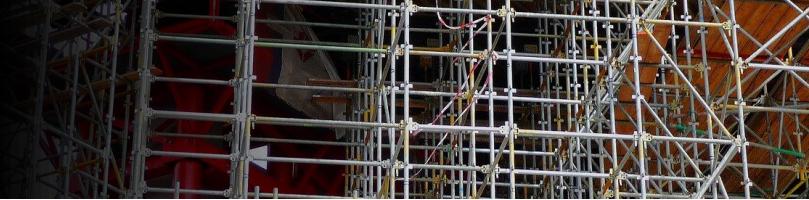
TRAIN AS BEFORE

# Transfer learning

Replace the output layer with a layer that produces the right number of outputs.

**Feature extraction:** Only calibrate the final layer, but "freeze" the rest of the network.

# Transfer learning



Replace the output layer with a layer that produces the right number of outputs.

**Feature extraction:** Only calibrate the final layer, but "freeze" the rest of the network.

```
model2 = torchvision.models.resnet18(pretrained=True) # You could also use resnet50 as above
```

```
# Freeze the network
```

```
for param in model2.parameters():
    param.requires_grad = False
```

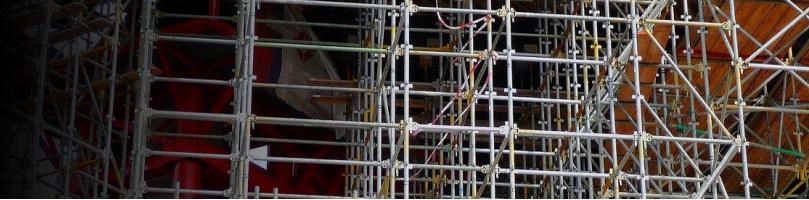
```
# Input features of the last layer of the original model
```

```
n_ftrs = model2.fc.in_features
```

```
n_classes = len(class_names)
```

```
model2.fc = nn.Linear(n_ftrs, n_classes)
```

# Transfer learning



Replace the output layer with a layer that produces the right number of outputs.

**Feature extraction:** Only calibrate the final layer, but "freeze" the rest of the network.

```
model2 = torchvision.models.resnet18(pretrained=True) # You could also use resnet50 as above
```

*# Freeze the network*

```
for param in model2.parameters():
    param.requires_grad = False
```

*# Input features of the last layer of the original model*

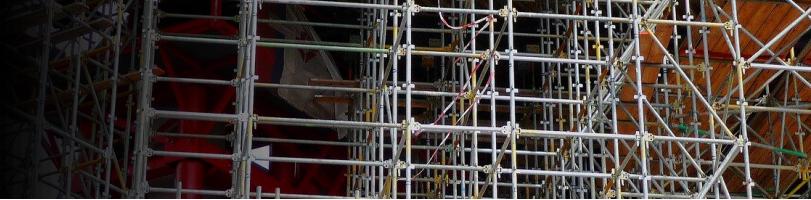
```
n_ftrs = model2.fc.in_features
```

```
n_classes = len(class_names)
```

```
model2.fc = nn.Linear(n_ftrs, n_classes)
```

TRAIN AS BEFORE

# Including GPUs



*# Check GPU availability*

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

Move the model to the GPUs

*# move to GPU*

```
model = model.to(device)
```

Move mini-batches to the GPUs

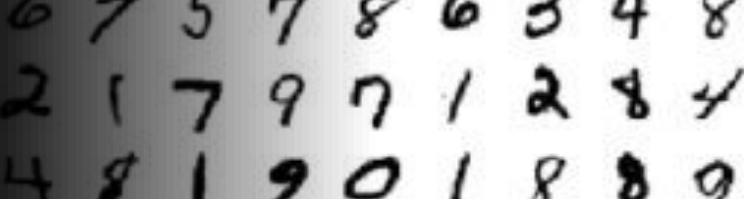
*# Modify training loop*

```
for epoch in range(num_epochs):
    for inputs, labels in trainloader:
        inputs = inputs.to(device)
        labels = labels.to(device)
```

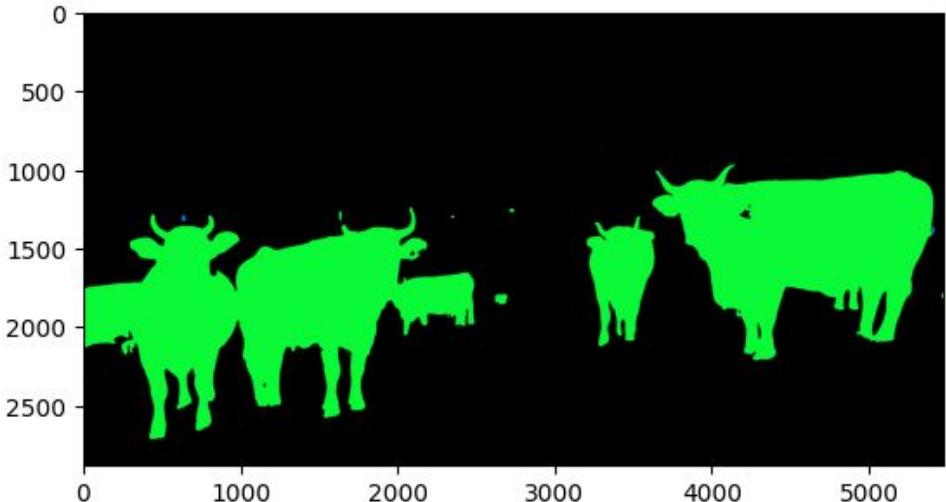
# Applications of CNNs

6	8	1	7	9	6	6	9
7	5	7	8	6	3	4	8
7	9	7	1	2	8	4	
7	1	9	0	1	8	8	9
6	1	8	6	4	1	5	6
5	9	2	6	5	8	1	9
2	2	2	2	3	4	4	8
3	3	8	0	7	3	8	5
4	6	4	6	0	2	4	
1	2	8	7	6	9	8	6

# Semantic segmentation

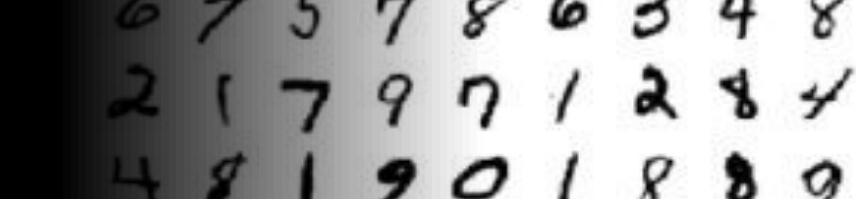
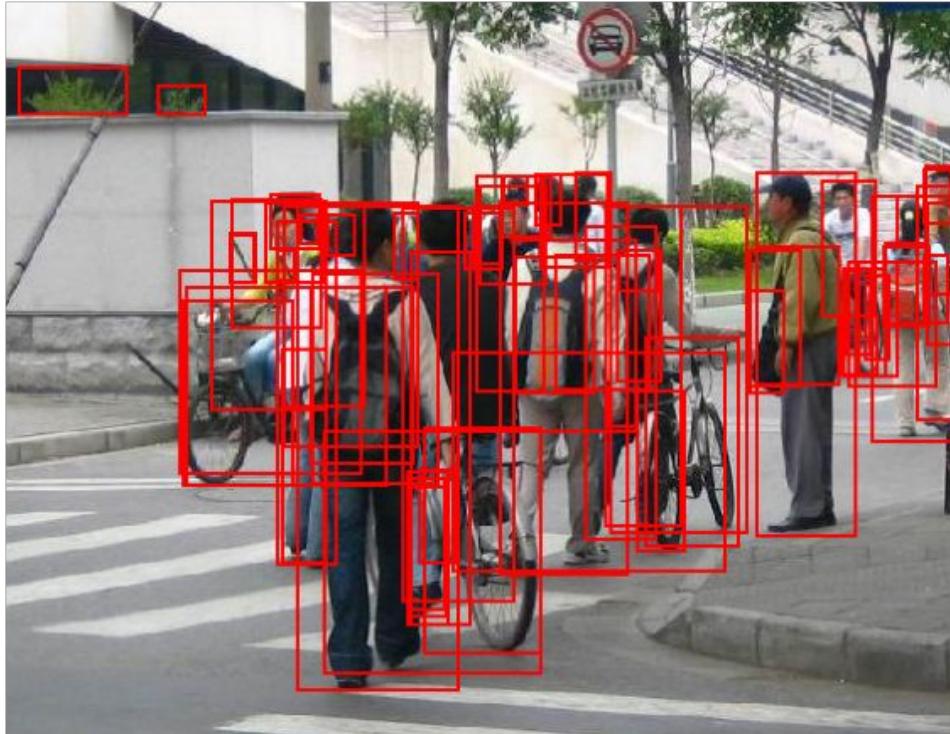


Use CNNs to assign a semantic label to every pixel.



```
model = torch.hub.load('pytorch/vision:v0.10.0', 'deeplabv3_resnet50', pretrained=True) # Pre-trained network (DeepLab3)
```

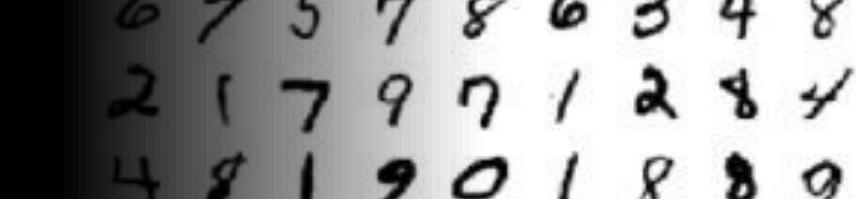
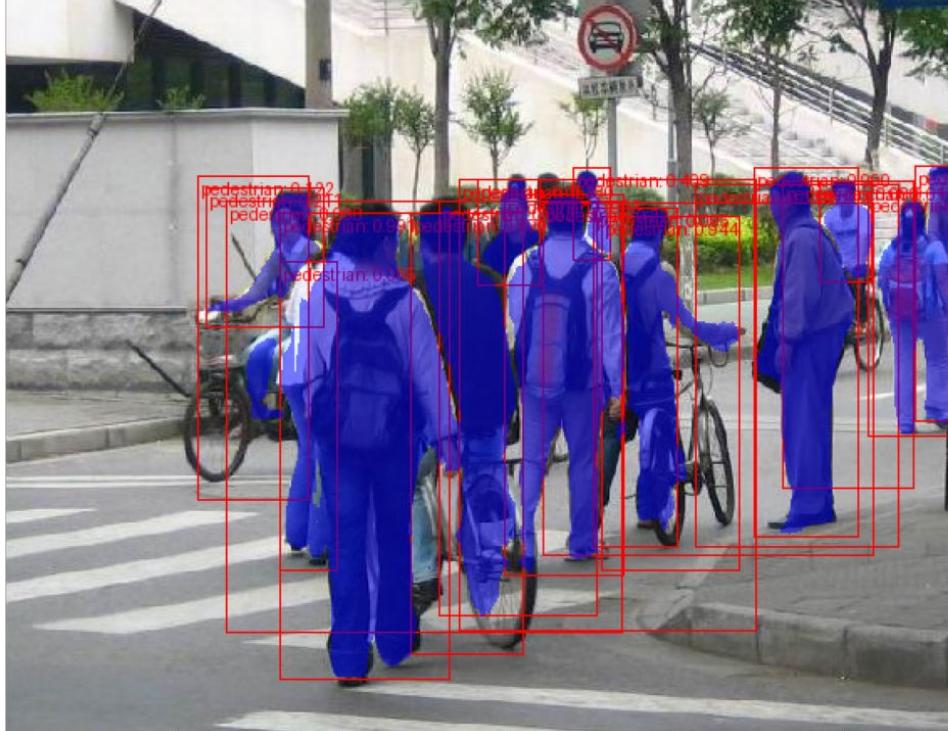
# Object detection



Use CNNs to detect and localise objects within an image.

Cf. PyTorch Tutorials on [pytorch.org](https://pytorch.org)

# Instance segmentation



Use CNNs to detect and localise objects within an image.

Cf. PyTorch Tutorials on [pytorch.org](https://pytorch.org)

Add instance segmentation to follow individuals.

# Image style transfer

Use CNNs to transfer the artistic style of one image to another.



+



=



Content

Style

New image

```
from torchvision.models import vgg19, VGG19_Weights # Pre-trained network
```

# Other examples



- CNNs are widely used for **face recognition** tasks.
- With Recurrent Neural Networks (RNNs), CNNs can **generate image captions**.
- CNNs are used in **generative models** like GANs (Generative Adversarial Networks) for generating realistic images.
- CNNs can be extended to **video data**.
- Super-Resolution: CNNs can be used for **enhancing the resolution** of images.

# Summary

Based on this lecture, you should better be able to

- 1) Explain basic concepts behind Convolutional Neural Networks (CNNs): convolutional and pooling layers). Outline the architecture of CNNs.
- 2) Translate these concepts into PyTorch.
- 3) Train and run neural networks on GPUs.
- 4) Describe the breadth of applications of CNNs.